

```
from google.colab import drive
drive.mount('/content/drive')
```

↻ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
import os
data_dir = "/content/drive/MyDrive/chest_xray 2"
```

```
!pip install tensorflow
```

```
!pip install opencv-python
```

↻ Requirement already satisfied: opencv-python in /usr/local/lib/python3.11/dist-packages (4.11.0.86)
Requirement already satisfied: numpy>=1.21.2 in /usr/local/lib/python3.11/dist-packages (from opencv-python) (2.0.2)

✓ import required library

```
import matplotlib.pyplot as plt
import seaborn as sns
import keras
from keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPool2D, Dense, Flatten, Dropout, BatchNormalization # Import MaxPooling2D from tensorflow.keras.layers
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
from keras.callbacks import ReduceLROnPlateau
import cv2
import os
import numpy as np
import pandas as pd
```

```
import cv2
import os
import numpy as np
```

```
labels = ['PNEUMONIA', 'NORMAL']
img_size = 150
```

```
def get_training_data(data_dir):
    data = []
```

```

for label in labels: # Use the global 'labels' variable
    path = os.path.join(data_dir, label)
    class_num = labels.index(label)
    for img in os.listdir(path):
        try:
            img_arr = cv2.imread(os.path.join(path, img), cv2.IMREAD_GRAYSCALE)
            # Ensure all images are resized to the same shape
            resized_arr = cv2.resize(img_arr, (img_size, img_size))
            data.append([resized_arr, class_num])
        except Exception as e:
            print(e)

# Convert the list of images to a NumPy array with a consistent shape
images = [item[0] for item in data]
# Change the variable name for extracted labels to avoid conflict
image_labels = [item[1] for item in data]

# Reshape images to have a consistent channel dimension
images = np.array(images)[: , : , :, np.newaxis]

return images, np.array(image_labels) # Return the new label variable

```

✓ LOADING THE DATASET

```

train = get_training_data("/content/drive/MyDrive/chest_xray 2/train")
test = get_training_data("/content/drive/MyDrive/chest_xray 2/test")
val = get_training_data("/content/drive/MyDrive/chest_xray 2/val")

```


✓ DATA VISUALIZATION & PREPROCESSING

```

l = []
for i in train[1]: # Iterate through the labels in train[1]
    if i == 0: # Check if the individual label is 0
        l.append("Pneumonia")
    else:
        l.append("Normal")
sns.set_style('darkgrid')
sns.countplot(l)

```


<Axes: xlabel='count'>





✓ previewing the image

```
plt.figure(figsize = (5,5))
plt.imshow(train[0][0], cmap='gray')
plt.title(labels[train[1][0]]) # Access the label using train[1][0]
```

```
plt.figure(figsize = (5,5))
plt.imshow(train[0][-1], cmap='gray') # Access the image using train[0][-1]
plt.title(labels[train[1][-1]]) # Access the label using train[1][-1]
```

 Text(0.5, 1.0, 'NORMAL')



```
x_train = train[0] # Access the images from train
y_train = train[1] # Access the labels from train

x_val = val[0] # Access the images from val
y_val = val[1] # Access the labels from val

x_test = test[0] # Access the images from test
y_test = test[1] # Access the labels from test

## normalize the data
x_train = np.array(x_train) / 255
x_val = np.array(x_val) / 255
x_test = np.array(x_test) / 255

## resize the data
x_train = x_train.reshape(-1, img_size, img_size, 1)
y_train = np.array(y_train)

x_val = x_val.reshape(-1, img_size, img_size, 1)
y_val = np.array(y_val)


x_test = x_test.reshape(-1, img_size, img_size, 1)
y_test = np.array(y_test)
```

✓ data Augmentation

```
datagen = ImageDataGenerator(
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    rotation_range = 30, # randomly rotate images in the range (degrees, 0 to 180)
    zoom_range = 0.2, # Randomly zoom image
    width_shift_range=0.1, # randomly shift images horizontally (fraction of total width)
    height_shift_range=0.1, # randomly shift images vertically (fraction of total height)
    horizontal_flip = True, # randomly flip images
    vertical_flip=False) # randomly flip images
```

✓ TRAIN THE MODEL

```
model = Sequential()  
model.add(Conv2D(32 , (3,3) , strides = 1 , padding = 'same' , activation = 'relu' , input_shape = (150,150,1)))  
model.add(BatchNormalization())  
model.add(MaxPool2D((2,2) , strides = 2 , padding = 'same'))  
model.add(Conv2D(64 , (3,3) , strides = 1 , padding = 'same' , activation = 'relu'))  
model.add(Dropout(0.1))  
model.add(BatchNormalization())  
model.add(MaxPool2D((2,2) , strides = 2 , padding = 'same'))  
model.add(Conv2D(64 , (3,3) , strides = 1 , padding = 'same' , activation = 'relu'))  
model.add(BatchNormalization())  
model.add(MaxPool2D((2,2) , strides = 2 , padding = 'same'))  
model.add(Conv2D(128 , (3,3) , strides = 1 , padding = 'same' , activation = 'relu'))  
model.add(Dropout(0.2))  
model.add(BatchNormalization())  
model.add(MaxPool2D((2,2) , strides = 2 , padding = 'same'))  
model.add(Conv2D(256 , (3,3) , strides = 1 , padding = 'same' , activation = 'relu'))  
model.add(Dropout(0.2))  
model.add(BatchNormalization())  
model.add(MaxPool2D((2,2) , strides = 2 , padding = 'same'))  
model.add(Flatten())  
model.add(Dense(units = 128 , activation = 'relu'))  
model.add(Dropout(0.2))  
model.add(Dense(units = 1 , activation = 'sigmoid'))  
model.compile(optimizer = "rmsprop" , loss = 'binary_crossentropy' , metrics = ['accuracy'])  
model.summary()
```

 /usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer that inherits from `Layer` without overriding the `input_shape` attribute.
 super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 150, 150, 32)	320
batch_normalization (BatchNormalization)	(None, 150, 150, 32)	128
max_pooling2d (MaxPooling2D)	(None, 75, 75, 32)	0
conv2d_1 (Conv2D)	(None, 75, 75, 64)	18,496
dropout (Dropout)	(None, 75, 75, 64)	0
batch_normalization_1 (BatchNormalization)	(None, 75, 75, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 38, 38, 64)	0
conv2d_2 (Conv2D)	(None, 38, 38, 64)	36,928
batch_normalization_2 (BatchNormalization)	(None, 38, 38, 64)	256
max_pooling2d_2 (MaxPooling2D)	(None, 19, 19, 64)	0
conv2d_3 (Conv2D)	(None, 19, 19, 128)	73,856
dropout_1 (Dropout)	(None, 19, 19, 128)	0
batch_normalization_3 (BatchNormalization)	(None, 19, 19, 128)	512
max_pooling2d_3 (MaxPooling2D)	(None, 10, 10, 128)	0
conv2d_4 (Conv2D)	(None, 10, 10, 256)	295,168
dropout_2 (Dropout)	(None, 10, 10, 256)	0
batch_normalization_4 (BatchNormalization)	(None, 10, 10, 256)	1,024
max_pooling2d_4 (MaxPooling2D)	(None, 5, 5, 256)	0
flatten (Flatten)	(None, 6400)	0
dense (Dense)	(None, 128)	819,328
dropout_3 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 1)	129

Total params: 1,246,401 (4.75 MB)

Trainable params: 1,245,313 (4.75 MB)

Non-trainable params: 1,088 (4.25 KB)

```
learning_rate_reduction = ReduceLRonPlateau(monitor='val_accuracy', patience = 2, verbose=1, factor=0.3, min_lr=0.000001)
```

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

```
history = model.fit(datagen.flow(x_train,y_train, batch_size = 32) ,epochs =25, validation_data = datagen.flow(x_val, y_val) ,callbacks = [learning_rate_reduction])
```

```
Epoch 1/25
163/163 ————— 475s 3s/step - accuracy: 0.9694 - loss: 0.0948 - val_accuracy: 0.5625 - val_loss: 1.5085 - learning_rate: 9.0000e-05
Epoch 2/25
163/163 ————— 462s 3s/step - accuracy: 0.9668 - loss: 0.1000 - val_accuracy: 0.5625 - val_loss: 0.9743 - learning_rate: 9.0000e-05
Epoch 3/25
163/163 ————— 467s 3s/step - accuracy: 0.9691 - loss: 0.0953 - val_accuracy: 1.0000 - val_loss: 0.1724 - learning_rate: 9.0000e-05
Epoch 4/25
163/163 ————— 461s 3s/step - accuracy: 0.9637 - loss: 0.0968 - val_accuracy: 0.6875 - val_loss: 0.7948 - learning_rate: 9.0000e-05
Epoch 5/25
163/163 ————— 0s 3s/step - accuracy: 0.9666 - loss: 0.0921
Epoch 5: ReduceLRonPlateau reducing learning rate to 2.700000040931627e-05.
163/163 ————— 467s 3s/step - accuracy: 0.9666 - loss: 0.0921 - val_accuracy: 0.5625 - val_loss: 2.0688 - learning_rate: 9.0000e-05
Epoch 6/25
163/163 ————— 468s 3s/step - accuracy: 0.9747 - loss: 0.0799 - val_accuracy: 0.5625 - val_loss: 1.8878 - learning_rate: 2.7000e-05
Epoch 7/25
163/163 ————— 0s 3s/step - accuracy: 0.9737 - loss: 0.0754
Epoch 7: ReduceLRonPlateau reducing learning rate to 8.100000013655517e-06.
163/163 ————— 510s 3s/step - accuracy: 0.9737 - loss: 0.0755 - val_accuracy: 0.6875 - val_loss: 0.8649 - learning_rate: 2.7000e-05
Epoch 8/25
163/163 ————— 469s 3s/step - accuracy: 0.9673 - loss: 0.0949 - val_accuracy: 0.5625 - val_loss: 1.0631 - learning_rate: 8.1000e-06
Epoch 9/25
163/163 ————— 0s 3s/step - accuracy: 0.9691 - loss: 0.0885
Epoch 9: ReduceLRonPlateau reducing learning rate to 2.429999949526973e-06.
163/163 ————— 470s 3s/step - accuracy: 0.9691 - loss: 0.0884 - val_accuracy: 0.6250 - val_loss: 1.3179 - learning_rate: 8.1000e-06
Epoch 10/25
163/163 ————— 470s 3s/step - accuracy: 0.9666 - loss: 0.0914 - val_accuracy: 0.6250 - val_loss: 1.4770 - learning_rate: 2.4300e-06
Epoch 11/25
163/163 ————— 0s 3s/step - accuracy: 0.9686 - loss: 0.0819
Epoch 11: ReduceLRonPlateau reducing learning rate to 1e-06.
163/163 ————— 468s 3s/step - accuracy: 0.9686 - loss: 0.0819 - val_accuracy: 0.6250 - val_loss: 1.4073 - learning_rate: 2.4300e-06
Epoch 12/25
163/163 ————— 463s 3s/step - accuracy: 0.9657 - loss: 0.0976 - val_accuracy: 0.6250 - val_loss: 1.1731 - learning_rate: 1.0000e-06
Epoch 13/25
163/163 ————— 504s 3s/step - accuracy: 0.9660 - loss: 0.0921 - val_accuracy: 0.6875 - val_loss: 0.9856 - learning_rate: 1.0000e-06
Epoch 14/25
163/163 ————— 457s 3s/step - accuracy: 0.9726 - loss: 0.0894 - val_accuracy: 0.7500 - val_loss: 0.7758 - learning_rate: 1.0000e-06
Epoch 15/25
163/163 ————— 452s 3s/step - accuracy: 0.9671 - loss: 0.0921 - val_accuracy: 0.5625 - val_loss: 1.2353 - learning_rate: 1.0000e-06
Epoch 16/25
163/163 ————— 454s 3s/step - accuracy: 0.9730 - loss: 0.0796 - val_accuracy: 0.6875 - val_loss: 0.9578 - learning_rate: 1.0000e-06
```


Epoch 17/25

163/163 ————— **465s** 3s/step - accuracy: 0.9731 - loss: 0.0794 - val_accuracy: 0.5625 - val_loss: 1.2103 - learning_rate: 1.0000e-06

Epoch 18/25

163/163 ————— **497s** 3s/step - accuracy: 0.9716 - loss: 0.0859 - val_accuracy: 0.6875 - val_loss: 0.6033 - learning_rate: 1.0000e-06

Epoch 19/25

163/163 ————— **511s** 3s/step - accuracy: 0.9727 - loss: 0.0782 - val_accuracy: 0.6250 - val_loss: 1.1148 - learning_rate: 1.0000e-06

Epoch 20/25

163/163 ————— **469s** 3s/step - accuracy: 0.9762 - loss: 0.0695 - val_accuracy: 0.6250 - val_loss: 0.9790 - learning_rate: 1.0000e-06

Epoch 21/25

163/163 ————— **467s** 3s/step - accuracy: 0.9698 - loss: 0.0772 - val_accuracy: 0.5625 - val_loss: 1.0537 - learning_rate: 1.0000e-06

Epoch 22/25

163/163 ————— **502s** 3s/step - accuracy: 0.9661 - loss: 0.0989 - val_accuracy: 0.6875 - val_loss: 1.0839 - learning_rate: 1.0000e-06

Epoch 23/25

163/163 ————— **468s** 3s/step - accuracy: 0.9701 - loss: 0.0858 - val_accuracy: 0.6250 - val_loss: 0.7873 - learning_rate: 1.0000e-06

Epoch 24/25

163/163 ————— **466s** 3s/step - accuracy: 0.9684 - loss: 0.0865 - val_accuracy: 0.6875 - val_loss: 1.0272 - learning_rate: 1.0000e-06

Epoch 25/25

163/163 ————— **471s** 3s/step - accuracy: 0.9709 - loss: 0.0825 - val accuracy: 0.8125 - val loss: 0.7226 - learning rate: 1.0000e-06Start coding or [generate](#) with AI.Start coding or [generate](#) with AI.Start coding or [generate](#) with AI.Start coding or [generate](#) with AI.Start coding or [generate](#) with AI.

```

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers

# Define a CNN Model with Dropout
# Changed input_shape to match the actual input data shape (150, 150, 1)
model = keras.Sequential([
    layers.Conv2D(32, (3,3), activation='relu', input_shape=(150, 150, 1)),
    layers.MaxPooling2D((2,2)),
    layers.Dropout(0.25), # Dropout after the first pooling layer

    layers.Conv2D(64, (3,3), activation='relu'),
    layers.MaxPooling2D((2,2)),
    layers.Dropout(0.25), # Dropout after the second pooling layer

    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.5), # Dropout before the final dense layer
    layers.Dense(1, activation='sigmoid') # Output layer for binary classification
])

```

```
# Compile the model
model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy'])
```

```
# Now, when you call model.fit(), make sure 'train' and 'val' data
# have the correct shape (num_samples, 150, 150, 1).
```

```
history = model.fit(train[0], train[1], validation_data=(val[0], val[1]), epochs=20) # Pass train[0] for images and train[1] for labels
```

```
Epoch 1/20
163/163 ————— 11s 23ms/step - accuracy: 0.7057 - loss: 105.0736 - val_accuracy: 0.6875 - val_loss: 0.8636
Epoch 2/20
163/163 ————— 4s 18ms/step - accuracy: 0.8727 - loss: 0.3102 - val_accuracy: 0.7500 - val_loss: 0.9732
Epoch 3/20
163/163 ————— 3s 18ms/step - accuracy: 0.8861 - loss: 0.2771 - val_accuracy: 0.7500 - val_loss: 1.2621
Epoch 4/20
163/163 ————— 5s 17ms/step - accuracy: 0.8995 - loss: 0.2554 - val_accuracy: 0.7500 - val_loss: 1.1577
Epoch 5/20
163/163 ————— 5s 18ms/step - accuracy: 0.9218 - loss: 0.2086 - val_accuracy: 0.7500 - val_loss: 1.2889
Epoch 6/20
163/163 ————— 3s 18ms/step - accuracy: 0.9159 - loss: 0.2133 - val_accuracy: 0.7500 - val_loss: 1.2214
Epoch 7/20
163/163 ————— 3s 17ms/step - accuracy: 0.9302 - loss: 0.1846 - val_accuracy: 0.6875 - val_loss: 1.2817
Epoch 8/20
163/163 ————— 3s 18ms/step - accuracy: 0.9402 - loss: 0.1537 - val_accuracy: 0.7500 - val_loss: 1.1582
Epoch 9/20
163/163 ————— 3s 17ms/step - accuracy: 0.9506 - loss: 0.1255 - val_accuracy: 0.7500 - val_loss: 1.7910
Epoch 10/20
163/163 ————— 5s 18ms/step - accuracy: 0.9542 - loss: 0.1339 - val_accuracy: 0.8125 - val_loss: 1.6373
Epoch 11/20
163/163 ————— 5s 18ms/step - accuracy: 0.9507 - loss: 0.1278 - val_accuracy: 0.8125 - val_loss: 2.1183
Epoch 12/20
163/163 ————— 5s 18ms/step - accuracy: 0.9727 - loss: 0.0802 - val_accuracy: 0.7500 - val_loss: 1.6384
Epoch 13/20
163/163 ————— 3s 18ms/step - accuracy: 0.9691 - loss: 0.0781 - val_accuracy: 0.8125 - val_loss: 1.8179
Epoch 14/20
163/163 ————— 3s 17ms/step - accuracy: 0.9748 - loss: 0.0714 - val_accuracy: 0.7500 - val_loss: 1.8844
Epoch 15/20
163/163 ————— 5s 18ms/step - accuracy: 0.9788 - loss: 0.0624 - val_accuracy: 0.8125 - val_loss: 1.9619
Epoch 16/20
163/163 ————— 5s 18ms/step - accuracy: 0.9790 - loss: 0.0632 - val_accuracy: 0.7500 - val_loss: 1.9582
Epoch 17/20
163/163 ————— 3s 18ms/step - accuracy: 0.9772 - loss: 0.0600 - val_accuracy: 0.8125 - val_loss: 2.5408
Epoch 18/20
163/163 ————— 3s 17ms/step - accuracy: 0.9845 - loss: 0.0436 - val_accuracy: 0.7500 - val_loss: 2.8651
Epoch 19/20
163/163 ————— 5s 19ms/step - accuracy: 0.9801 - loss: 0.0531 - val_accuracy: 0.7500 - val_loss: 3.4313
Epoch 20/20
163/163 ————— 3s 18ms/step - accuracy: 0.9823 - loss: 0.0500 - val_accuracy: 0.8125 - val_loss: 2.9960
```

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.


```
import os
train_normal = len(os.listdir('/content/drive/MyDrive/chest_xray 2/train/NORMAL'))
train_pneumonia = len(os.listdir('/content/drive/MyDrive/chest_xray 2/train/PNEUMONIA'))
print(f"Normal: {train_normal}, Pneumonia: {train_pneumonia}")
```

🔗 Normal: 1341, Pneumonia: 3875

```
## import required library
import os
import matplotlib.pyplot as plt
import seaborn as sns
print(f"Normal: {train_normal}, Pneumonia: {train_pneumonia}")

# Plot class distribution
sns.barplot(x=['Normal', 'Pneumonia'], y=[train_normal, train_pneumonia])
plt.title("Class Distribution in Training Data")
plt.ylabel("Number of Images")
plt.show()
```

↔ Normal: 1341, Pneumonia: 3875



```
## we need to balanced the data by using data augumentation
from tensorflow.keras.preprocessing.image import ImageDataGenerator

# Data Augmentation
train_datagen = ImageDataGenerator(
    rescale=1./255, # Normalize pixel values
    rotation_range=20, # Randomly rotate images
    width_shift_range=0.2, # Horizontal shift
    height_shift_range=0.2, # Vertical shift
    shear_range=0.2, # Shear transformation
    zoom_range=0.2, # Random zoom
    horizontal_flip=True, # Flip images
    fill_mode='nearest', # Fill missing pixels
    validation_split=0.2 # Split for validation
)

test_datagen = ImageDataGenerator(rescale=1./255) # No augmentation for test set

# Load Train Dataset
train_generator = train_datagen.flow_from_directory(
    os.path.join(data_dir, 'train'),
    target_size=(150, 150), # Resize images
    batch_size=32,
    class_mode='binary', # Binary classification
```


```

    subset='training' # Training subset
)

# Load Validation Dataset
validation_generator = train_datagen.flow_from_directory(
    os.path.join(data_dir, 'train'),
    target_size=(150, 150),
    batch_size=32,
    class_mode='binary',
    subset='validation' # Validation subset
)

# Load Test Dataset
test_generator = test_datagen.flow_from_directory(
    os.path.join(data_dir, 'test'),
    target_size=(150, 150),
    batch_size=32,
    class_mode='binary'
)

```

 Found 4173 images belonging to 2 classes.
 Found 1043 images belonging to 2 classes.
 Found 624 images belonging to 2 classes.

```

from sklearn.utils.class_weight import compute_class_weight
import numpy as np

```

```

# Define class labels
class_labels = ['NORMAL', 'PNEUMONIA']
class_counts = [train_normal, train_pneumonia]

```

```


# Compute class weights
weights = compute_class_weight(class_weight="balanced", classes=np.array([0, 1]), y=[0]*train_normal + [1]*train_pneumonia)
class_weights = {0: weights[0], 1: weights[1]}

```

```

print(f"Class Weights: {class_weights}")

```

 Class Weights: {0: np.float64(1.9448173005219984), 1: np.float64(0.6730322580645162)}

```

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout

```

```

# CNN Model Architecture
model = Sequential([
    Conv2D(32, (3,3), activation='relu', input_shape=(150, 150, 3)),
    MaxPooling2D(2,2),

```

```

Conv2D(64, (3,3), activation='relu'),
MaxPooling2D(2,2),

Conv2D(128, (3,3), activation='relu'),
MaxPooling2D(2,2),

Flatten(),
Dense(128, activation='relu'),
Dropout(0.5),
Dense(1, activation='sigmoid') # Binary Classification (Normal/Pneumonia)
])

# Compile Model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# Train Model with Class Weights
history = model.fit(train_generator,
                    validation_data=validation_generator,
                    epochs=10,
                    class_weight=class_weights) # Apply class weights

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__(self._warn_if_super_not_called())
Epoch 1/10
131/131 ————— 95s 692ms/step - accuracy: 0.4555 - loss: 0.7328 - val_accuracy: 0.7239 - val_loss: 0.5847
Epoch 2/10
131/131 ————— 83s 638ms/step - accuracy: 0.8123 - loss: 0.4124 - val_accuracy: 0.8102 - val_loss: 0.3906
Epoch 3/10
131/131 ————— 88s 675ms/step - accuracy: 0.8460 - loss: 0.3513 - val_accuracy: 0.8159 - val_loss: 0.3866
Epoch 4/10
131/131 ————— 84s 642ms/step - accuracy: 0.8395 - loss: 0.3697 - val_accuracy: 0.8293 - val_loss: 0.3933
Epoch 5/10
131/131 ————— 86s 661ms/step - accuracy: 0.8664 - loss: 0.3099 - val_accuracy: 0.8207 - val_loss: 0.4286
Epoch 6/10
131/131 ————— 85s 645ms/step - accuracy: 0.8817 - loss: 0.3001 - val_accuracy: 0.8428 - val_loss: 0.3584
Epoch 7/10
131/131 ————— 84s 646ms/step - accuracy: 0.8845 - loss: 0.2771 - val_accuracy: 0.8821 - val_loss: 0.2613
Epoch 8/10
131/131 ————— 88s 676ms/step - accuracy: 0.8838 - loss: 0.2696 - val_accuracy: 0.8830 - val_loss: 0.2784
Epoch 9/10
131/131 ————— 84s 641ms/step - accuracy: 0.8876 - loss: 0.2575 - val_accuracy: 0.8974 - val_loss: 0.2486
Epoch 10/10
131/131 ————— 85s 652ms/step - accuracy: 0.8904 - loss: 0.2852 - val_accuracy: 0.8869 - val_loss: 0.2621


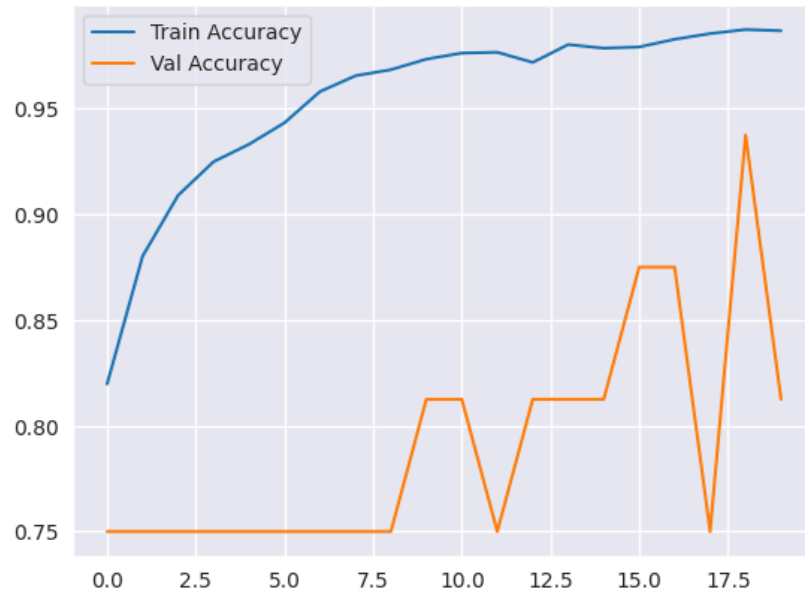
import matplotlib.pyplot as plt

# Plot training & validation accuracy
plt.plot(history.history['accuracy'], label='Train Accuracy')

```

```
plt.plot(history.history['val_accuracy'], label='Val Accuracy')
plt.legend()
plt.show()

# Plot training & validation loss
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.legend()
plt.show()
```



```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
train_datagen = ImageDataGenerator(  
    rescale=1./255,  
    rotation_range=20,
```



```

width_shift_range=0.2,
height_shift_range=0.2,
shear_range=0.2,
zoom_range=0.2,
horizontal_flip=True,
fill_mode='nearest'
)

from tensorflow.keras.callbacks import EarlyStopping

early_stop = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)

history = model.fit(train_generator,
                    validation_data=validation_generator,
                    epochs=20,
                    class_weight=class_weights,
                    callbacks=[early_stop])

➡ Epoch 1/20
131/131 ————— 87s 663ms/step - accuracy: 0.9047 - loss: 0.2143 - val_accuracy: 0.8658 - val_loss: 0.3143
Epoch 2/20
131/131 ————— 84s 644ms/step - accuracy: 0.8928 - loss: 0.2502 - val_accuracy: 0.8878 - val_loss: 0.2675
Epoch 3/20
131/131 ————— 88s 672ms/step - accuracy: 0.8992 - loss: 0.2296 - val_accuracy: 0.9070 - val_loss: 0.2154
Epoch 4/20
131/131 ————— 84s 641ms/step - accuracy: 0.8953 - loss: 0.2447 - val_accuracy: 0.8629 - val_loss: 0.3196
Epoch 5/20
131/131 ————— 84s 646ms/step - accuracy: 0.8981 - loss: 0.2331 - val_accuracy: 0.8840 - val_loss: 0.2547
Epoch 6/20
131/131 ————— 84s 641ms/step - accuracy: 0.9095 - loss: 0.2167 - val_accuracy: 0.9137 - val_loss: 0.2224

from google.colab import drive
drive.mount('/content/drive')

# Save model to Google Drive
model.save('/content/drive/My Drive/pneumonia_model.h5')

➡ WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend saving your model as a SavedModel using `model.save('path/to/saved_model')`.
Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).


class_weights = {0: 1.0, 1: 3.0} # Increase PNEUMONIA weight

from tensorflow.keras.models import load_model

# Load the trained model

```

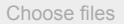
```
model = load_model('/content/drive/MyDrive/pneumonia_model.h5')
```

 WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile_metrics` will be empty until you train or evaluate the mc

✓ check the model on new data

```
from google.colab import files
```

```
# Upload an image
uploaded = files.upload()
```

 No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving IM-0025-0001.jpeg to IM-0025-0001 (1).jpeg

```
class_weights = {0: 1.0, 1: 3.0} # Increase PNEUMONIA weight
```

```
import numpy as np
import cv2
from tensorflow.keras.preprocessing import image
```

```
# Load and preprocess the uploaded image
img_path = list(uploaded.keys())[0] # Get uploaded filename
# Resize to match the input shape your model was trained on (e.g., 150x150)
img = image.load_img(img_path, target_size=(150, 150))
```

```
# Convert image to array and normalize
img_array = image.img_to_array(img) / 255.0
img_array = np.expand_dims(img_array, axis=0) # Add batch dimension
```

```
# Make prediction
prediction = model.predict(img_array)
```

```
# Get class labels (assuming you used 0 for 'NORMAL' and 1 for 'PNEUMONIA')
class_labels = ['NORMAL', 'PNEUMONIA']
predicted_class = class_labels[np.argmax(prediction)]
```

```
# Print the result
print(f"The model predicts: {predicted_class}")
```

 1/1 **0s** 41ms/step
The model predicts: NORMAL

```
img = image.load_img(img_path, target_size=(224, 224)) # Match model input size
img_array = image.img_to_array(img) / 255.0 # Normalize (important!)
```